CASE STUDY REPORT

TeleLens: Telecom Feedback Intelligence

Word2Vec, Sentiment Classification, Text Summarization and Factoid Question Answering in One Evaluated System

---

Artificial Intelligence / Natural Language Processing

Prepared by:	Bharat Chandra Bodapati
Student ID:	2023005380
Study domain:	Telecom customer feedback
Primary application:	Complaint triage, opinion digests and evidence-backed answers
Report date:	1 November 2026
Source code:	github.com/BharatChandra-sys/NLP-case-study
Submitted to:	Radha Karampudi, Assistant Professor Department of Computer Science and Systems Engineering GITAM School of Computer Science and Engineering, Hyderabad

Research integrity note. Every score in this report is generated from the result files in the linked repository, produced by notebooks executed on a Kaggle T4 GPU. Word pairs, questions and thresholds used for evaluation are separated from those used for fitting, and every change made after seeing test results is disclosed in Section 10. Results that were neutral or negative are reported alongside the positive ones. The factoid gold set is a single-annotator set and has not been independently validated.

Contents

Executive Summary **1**

1 Problem Statement **1**

1.1 Background and real-world need 1

1.2 Problem statement 1

1.3 Why the problem is hard 1

1.4 Objectives 2

2 Data and Experimental Protocol **2**

2.1 Datasets 2

2.2 Evaluation protocol and leakage control 2

3 System Architecture **3**

4 Case Study 1: Word2Vec **3**

4.1 Theory 3

4.2 Original contribution: Polarity-Aware Word2Vec (PA-W2V) 3

4.3 Implementation 4

4.4 Results 4

5 Case Study 2: Sentiment Classification **6**

5.1 Methods 6

5.2 Results 7

5.3 Original contribution: selective triage 8

5.4 Robustness 9

6 Case Study 3: Text Summarization **9**

6.1 Original contribution: PAB-MMR 9

6.2 Recovering sentences from unpunctuated text 9

6.3 Evaluation against human references 10

6.4 Evaluation on telecom digests 11

7 Case Study 4: Factoid Question Answering **12**

7.1 Pipeline 12

7.2 Why a new gold set was needed 12

7.3 Public benchmark 13

7.4 Retrieval 13

7.5 End-to-end results 13

7.6 Error analysis 15

8 Integrated System **15**

9 Comparison with Supplied References **16**

10 Threats to Validity and Disclosed Limitations **16**

11 Conclusion and Future Work **17**

References **17**

A Evidence Index **19**

B Execution Record **19**

C Reproducibility **19**

D Selected Implementation Fragment **21**

Executive Summary

This report extends an earlier telecom feedback triage prototype, which measured a TF–IDF classifier and described Word2Vec, into a single system covering all four assigned topics. Each topic is treated as a full case study with theory, implementation, a measured baseline comparison and an original contribution. The four contributions share one theme: *a language system should represent opposite meanings as different, preserve the balance of what customers said, and decline to decide when the evidence is weak.*

Table 1: Assignment coverage and headline evidence.

Assigned topic	Original contribution	Headline measured result	Evidence
Case Study on Word2Vec (Sec. 4)	PA-W2V: negation binding and polarity counter-fitting with labels nested in folds	Held-out antonym cosine 0.683 → 0.373; neighbour purity 0.773 → 0.839	`Notebook 01`
Case Study on Sentiment Classification (Sec. 5)	Calibrated selective triage with cross-fitted deferral thresholds	Macro-F1 0.916 (n = 9,874); 5% risk target met at 4.93% with 91% coverage	`Notebook 02`
Case Study on Text Summarization (Sec. 6)	PAB-MMR: polarity- and aspect-balanced selection; punctuation recovery for unpunctuated text	Polarity gap 0.059 → 0.019; aspect recall 0.680 → 0.850; ROUGE statistically tied	`Notebook 03`
Case Study on Factoid QA (Sec. 7)	Answer-type gating and calibrated abstention over complaints	Overall accuracy 0.477 → 0.602; answered EM 0.755; 80% of unanswerable rejected	`Notebook 04`

1 Problem Statement

1.1 Background and real-world need

Telecom providers receive short, informal feedback about broadband reliability, mobile coverage, billing, installation and support. Support teams must identify negative experiences quickly, understand what customers are saying in aggregate, and look up specific facts inside historical complaints. These are four different language tasks, and in practice they are usually built as disconnected tools. This project builds them as one system, **TeleLens**, over real public telecom reviews.

1.2 Problem statement

How can one NLP system use distributional word representations to classify telecom feedback, summarize it without distorting its sentiment, and answer factual questions with evidence, while measuring its failures and declining to decide when its evidence is weak?

1.3 Why the problem is hard

- **Antonyms share contexts.** "Stable" and "unstable" occur in near-identical sentences, so Word2Vec places them close together. The earlier report showed TF–IDF cosine of 0.872 between "the service is reliable" and "the service is not reliable".
- **Summaries can distort.** A summary that selects central sentences can silently turn a mixed set of reviews into an all-negative or all-positive one.
- **The main corpus has no punctuation or digits.** Sentence-based summarization cannot be applied directly, and numeric facts have been erased.
- **Wrong answers are costly.** A question answering system that always answers will invent a fact when the corpus does not contain one.

1.4 Objectives

1. Implement skip-gram Word2Vec from scratch, validate it against library implementations and improve its handling of opposite meanings.
2. Compare sentiment models under leakage-controlled evaluation and turn the best into a calibrated triage policy.
3. Build a summarizer that preserves sentiment balance and aspect coverage, and evaluate it against human references.
4. Build a factoid question answering pipeline with evidence and abstention, and evaluate it on a public benchmark and a domain gold set.
5. Integrate all four into a single working system and publish code, notebooks and results.

2 Data and Experimental Protocol

2.1 Datasets

Table 2: Datasets, role and provenance.

Dataset	Size used	Role and source
WEDataset `Data10KMixed`	9,874 reviews after removing blank rows and 13 conflicting-label groups; 617,573 tokens	Word2Vec, sentiment, telecom digests. HeshamAsem/WEDataset, linked from reference [1]
WEDataset `Data1KMixed`	999 reviews	Continuity with the earlier report
Comcast ConsumerAffairs	5,629 complaints; 34,054 passages	Factoid QA corpus; cross-domain sentiment test. Kaggle
FewSum Amazon gold	60 products, 8 reviews each, 3 human summaries	Summarization references: 40 dev, 20 test. abrazinskas/FewSum
SQuAD v1.1 dev	Seeded sample of 2,000 questions	Public reader benchmark. SQuAD
SimLex-999	999 pairs; about 250 in vocabulary	External similarity benchmark. SimLex-999
TeleLens factoid gold	102 answerable, 30 unanswerable; 44 dev, 88 test	Domain QA evaluation. `data/gold/qa_gold.jsonl`

The WEDataset repository declares no licence, so the raw file is downloaded at run time and is not redistributed in the project repository. The 1,000-row and 10,000-row files share only part of their text (15.4% of the 1K reviews occur in the 10K file) and must not be confused with the 367,000-review corpus described in reference [1].

2.2 Evaluation protocol and leakage control

- **Grouped five-fold validation.** Identical review texts stay in one fold (`StratifiedGroupKFold`, seed 42).
- **Fold-nested fitting.** Vocabulary, IDF, embeddings, polarity scores and classifier calibration are fitted on the training fold only.
- **Held-out word pairs.** Word pairs used to evaluate embeddings are excluded from counter-fitting constraints.
- **Cross-fitted thresholds.** A deferral threshold applied to one fold is chosen from the other four.
- **Development and test separation.** Summarizer weights, QA configuration and the abstention threshold are selected on development data only.

- **Paired statistics.** Bootstrap confidence intervals, exact McNemar tests and Wilcoxon signed-rank tests support each comparison.

3 System Architecture

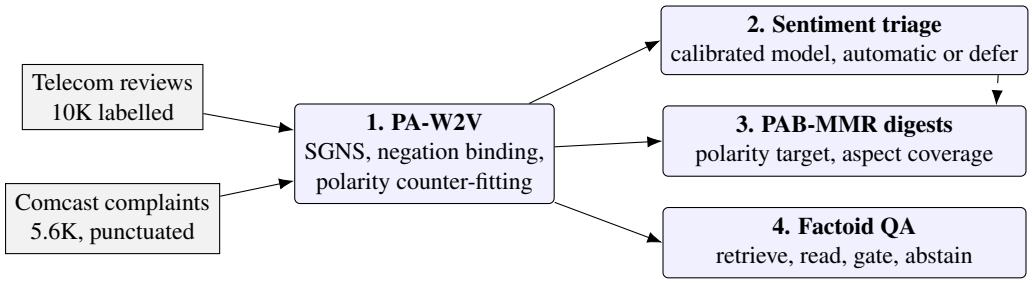


Figure 1: TeleLens architecture. Embeddings feed the downstream tasks; sentence polarity from the sentiment model feeds the summarizer (dashed).

Core logic is in `src/telelens`, one module per task. Notebooks in `notebooks` run the experiments, and `pipeline.py` exposes three operations: `triage`, `digest` and `ask`. Heavy computation ran on Kaggle (Tesla T4 GPU); the local machine holds code and pulled outputs.

4 Case Study 1: Word2Vec

4.1 Theory

Skip-gram with negative sampling (SGNS) [5,6] learns a target vector $\mathbf{v}_w$ and context vector $\mathbf{u}_c$. For an observed pair (w,c) and K noise words drawn from $P_n(w) \propto \text{count}(w)^{3/4}$ it minimises

$$\mathscr{L}_{\text{SGNS}} = -\log \sigma(\mathbf{u}_c^\top \mathbf{v}_w) - \sum_{i=1}^{K} \log \sigma(-\mathbf{u}_{n_i}^\top \mathbf{v}_w). \tag{1}$$

Using $\sigma'(x) = \sigma(x)\sigma(-x)$ gives the gradient

$$\frac{\partial \mathscr{L}_{\text{SGNS}}}{\partial \mathbf{v}_w} = -\sigma(-\mathbf{u}_c^\top \mathbf{v}_w)\,\mathbf{u}_c + \sum_{i=1}^{K} \sigma(\mathbf{u}_{n_i}^\top \mathbf{v}_w)\,\mathbf{u}_{n_i}, \tag{2}$$

which pulls $\mathbf{v}_w$ towards its observed contexts and pushes it from noise words. This explains the weakness studied here: if two words share contexts, the same $\mathbf{u}_c$ pull them to the same place, *whether they are synonyms or antonyms*. FastText [7] adds character n-grams but does not remove the effect and, because "stable" and "unstable" share character n-grams, can worsen it.

4.2 Original contribution: Polarity-Aware Word2Vec (PA-W2V)

PA-W2V adds two steps around SGNS, requiring no external lexicon.

Step 1: negation binding. Before training, a negator and the next content word become one token: "not very reliable" becomes `not_reliable`. This extends the negation marking of Pang et al. [11] to embedding training and gives negated meaning its own vector.

Step 2: polarity counter-fitting. Word polarity is estimated by a log-odds ratio with an informative Dirichlet prior [10], using only training-fold labels:

$$z_w = \frac{\log \frac{y_w^+ + a}{n^+ + a_0 - y_w^+ - a} - \log \frac{y_w^- + a}{n^- + a_0 - y_w^- - a}}{\sqrt{\frac{1}{y_w^+ + a} + \frac{1}{y_w^- + a}}}. \tag{3}$$

Among words with $|z_w| \geq 2.58$, each word's ten nearest polar neighbours generate an *attract* set A (same sign) and a *repel* set R (opposite sign). R also receives lexical antonyms found automatically in the vocabulary (w and `not_`w; w and un-, in-, dis-, im-, ir-w). Extending counter-fitting [8] and retrofitting [9], vectors are then adjusted by minimising

$$\mathscr{L}_{\mathrm{PA}} = \alpha \overline{\|\hat{\mathbf{v}}_w - \mathbf{v}_w^0\|^2} + \rho \left.\overline{\left[\cos(\hat{\mathbf{v}}_a, \hat{\mathbf{v}}_b)\right]_+}\right|_R + \kappa \left.\overline{\left[\mu - \cos(\hat{\mathbf{v}}_a, \hat{\mathbf{v}}_b)\right]_+}\right|_A, \tag{4}$$

where the first term preserves the original space, with $\alpha = 1$, $\rho = 1$, $\kappa = 0.3$ and $\mu = 0.6$. Because cosine lies in $[-1, 1]$ the repel term is zero exactly when a pair is orthogonal or opposed, so the loss stops pushing a pair once it is separated. The original method relies on external lexicons [8]; here constraints are derived from the corpus and its labels inside each fold.

4.3 Implementation

SGNS is written from scratch in PyTorch (`embeddings.py`): 200 dimensions, window 5, 10 negatives, 5 epochs, minimum count 5, subsampling 10^{-3}, SparseAdam. Gensim Word2Vec and FastText with matching settings provide library baselines. Counter-fitting weights were selected on an inner validation split (Table 3); validation macro-F1 varied little, so this choice has limited influence.

Table 3: Counter-fitting grid on the inner validation split of fold 0. The top row became the default.

Repel	Attract	Anchor	Val. macro-F1	Ant. cos	Syn. cos
1.0	**0.3**	**1.0**	**0.8880**	**0.489**	**0.743**
2.0	0.3	1.0	0.8868	0.488	0.738
2.0	0.6	2.0	0.8867	0.488	0.743
0.5	0.3	4.0	0.8861	0.537	0.786
1.0	0.1	4.0	0.8848	0.501	0.747
1.0	0.3	2.0	0.8836	0.491	0.745

4.4 Results

Table 4 uses 16 synonym, 16 antonym and 10 negation pairs that were excluded from the constraints. Purity is the share of a word's ten nearest polar neighbours with the same polarity, with polarity calculated from the *held-out* fold, so the labels used for fitting cannot inflate it.

Table 4: Word2Vec evaluation across five grouped folds (mean $\pm$ s.d.). Evaluation word pairs were never used as constraints.

Model	Vocab	Syn. cos	Ant. cos	Neg. cos	Syn/Ant AUC	Purity	SimLex ρ
SGNS (scratch)	3758	0.833 ± 0.007	0.695 ± 0.011	–	0.817 ± 0.039	0.774 ± 0.003	0.107 ± 0.020
gensim Word2Vec	3758	0.728 ± 0.002	0.645 ± 0.011	–	0.691 ± 0.023	0.706 ± 0.005	0.067 ± 0.027
FastText	3758	0.728 ± 0.005	0.732 ± 0.004	–	0.539 ± 0.010	0.709 ± 0.009	0.148 ± 0.011
SGNS + negation	4050	0.828 ± 0.008	0.683 ± 0.015	0.688 ± 0.030	0.823 ± 0.046	0.773 ± 0.012	0.099 ± 0.012
PA-W2V	**4050**	$\mathbf{0.742 \pm 0.015}$	$\mathbf{0.373 \pm 0.017}$	$\mathbf{0.559 \pm 0.061}$	$\mathbf{0.863 \pm 0.020}$	$\mathbf{0.839 \pm 0.021}$	$\mathbf{0.096 \pm 0.019}$

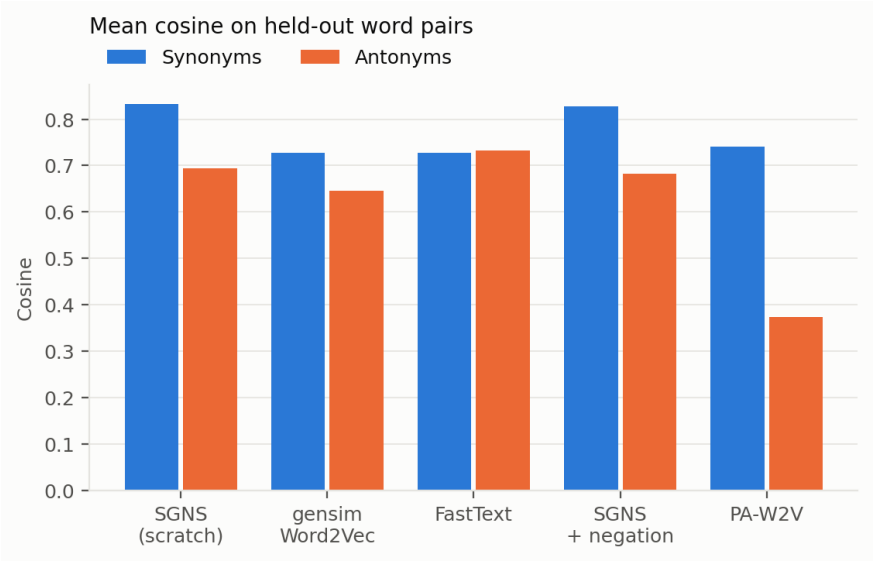


Figure 2: PA-W2V separates held-out antonyms from synonyms. FastText places antonyms as close as synonyms.

- **Opposites separate.** Mean antonym cosine falls from 0.683 to 0.373, and word–negation cosine from 0.688 to 0.559. The synonym-versus-antonym AUC rises from 0.823 to 0.863, while neighbour purity rises from 0.773 to 0.839.
- **There is a cost.** Synonym cosine also falls, from 0.828 to 0.742, so the space is less tightly clustered overall.
- **The from-scratch model is valid.** It matches or exceeds gensim on these checks. Its loss decreases monotonically by epoch (`sgns_loss.json`).
- **No change on SimLex-999 [27].** Spearman ρ stays near 0.10 for all models. SimLex is general English, and the corpus is small and domain-specific, so this benchmark shows no gain.
- **Pooling dilutes the effect.** Table 5 shows sentence cosine remains high after mean pooling, as most words are shared. The earlier report's concern is reduced at the word level but not solved at sentence level.

Table 5: Sentence-level cosine after mean pooling. Averaging dilutes the word-level gain.

Sentence pair	SGNS	FastText	PA-W2V
the service is reliable / the service is not reliable	0.976	0.972	0.955
the connection is stable / the connection is unstable	0.990	0.994	0.957
the engineer was helpful / the engineer was not helpful	0.971	0.967	0.946
my problem was fixed / my problem was not fixed	0.975	0.966	0.963
broadband is slow and keeps disconnecting / internet connection is sluggish and drops	0.962	0.937	0.928

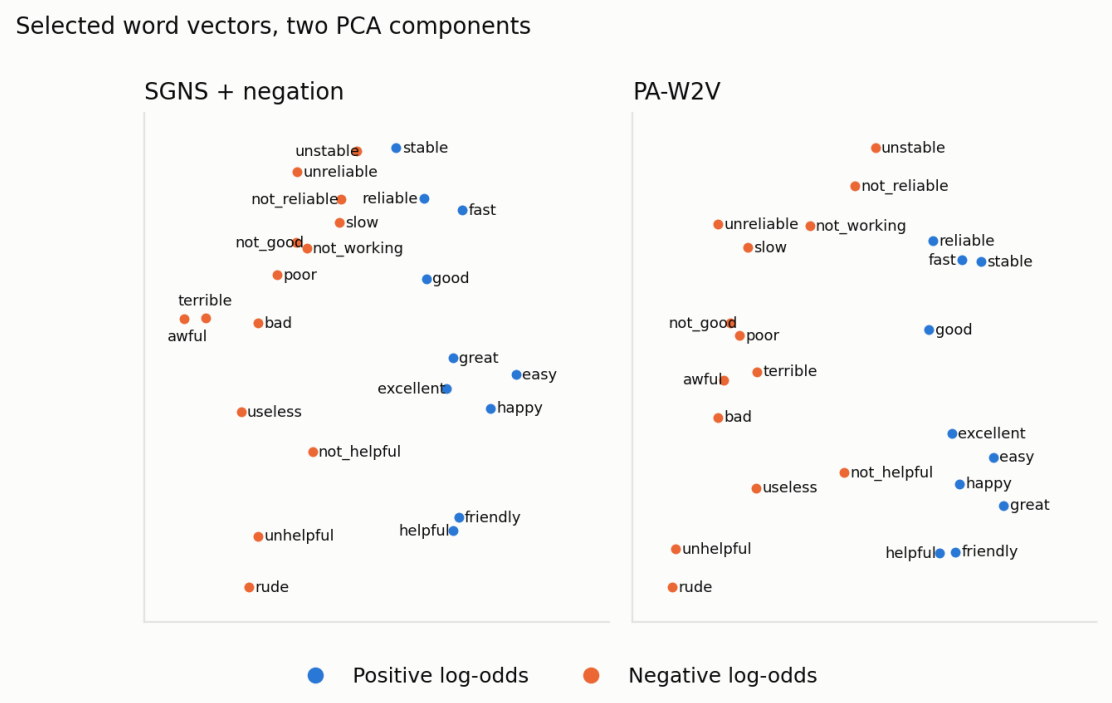


Figure 3: PCA of selected vectors. A two-dimensional projection discards information and is illustrative only; the quantitative evidence is in Table 4.

Table 6: Nearest neighbours on the whole-corpus model.

Word	Model	Top six neighbours
reliable	SGNS + negation	solid, stable, consistent, consistently, trustworthy, not_complaints
reliable	PA-W2V	consistently, trustworthy, lacking, rarely, importantly, mega
unreliable	SGNS + negation	disconnecting, reliability, dropouts, patchy, zoom, freezing
unreliable	PA-W2V	keeps, dropping, internet, disconnecting, awful, slow
stable	SGNS + negation	consistent, solid, rarely, reliable, slower, reliability
stable	PA-W2V	solid, competitors, pricey, rare, reliable, preferred
helpful	SGNS + negation	polite, patient, knowledgeable, professional, pleasant, courteous
helpful	PA-W2V	patient, polite, knowledgeable, courteous, knowledgable, kyle
bill	SGNS + negation	final, paper, payment, fees, declined, unpaid
bill	PA-W2V	final, paper, declined, fees, termination, advance

Table 6 also shows limits: *lacking* and *pricey* appear near positive words, reflecting co-occurrence rather than meaning. Small-corpus embeddings remain noisy.

5 Case Study 2: Sentiment Classification

5.1 Methods

Reviews are labelled Bad or Good from star ratings. All models use the same grouped folds. Linear models use logistic regression with Platt calibration [14] fitted inside the training fold. DistilBERT [12] is fine-tuned for two epochs (learning rate 2×10^{-5}, 256 tokens) and calibrated by temperature scaling [13] on a 10% validation split; fitted temperatures of 1.09–1.29 confirm mild overconfidence before calibration. An untrained SST-2 DistilBERT measures what domain training adds.

5.2 Results

Table 7: Out-of-fold sentiment results on 9,874 reviews (grouped five-fold).

Model	Accuracy	95% CI	Macro-F1	Bad recall	Bad prec.	ECE
DistilBERT fine-tuned	**0.916**	**0.911–0.921**	**0.916**	**0.910**	**0.922**	**0.024**
TF-IDF LR	0.897	0.891–0.903	0.897	0.920	0.881	0.013
TF-IDF + PA-W2V	0.896	0.890–0.902	0.895	0.922	0.877	0.019
TF-IDF + SGNS	0.891	0.885–0.897	0.891	0.918	0.872	0.024
PA-W2V mean LR	0.885	0.879–0.892	0.885	0.915	0.865	0.032
SGNS+neg mean LR	0.878	0.871–0.884	0.878	0.904	0.861	0.028
SGNS mean LR	0.871	0.865–0.878	0.871	0.902	0.852	0.031
FastText mean LR	0.857	0.850–0.864	0.857	0.895	0.834	0.036
DistilBERT SST-2 (no training)	0.823	0.815–0.831	0.819	0.971	0.751	0.159

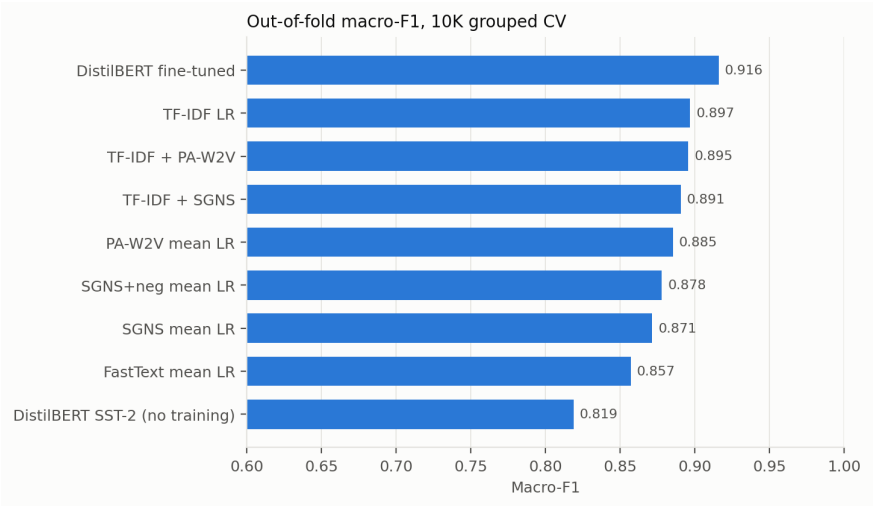


Figure 4: Out-of-fold macro-F1. The axis starts at 0.6 to show differences between models.

Table 8: Paired comparisons: bootstrap macro-F1 difference and exact McNemar test.

Comparison (A vs B)	F1 diff	95% CI	McNemar p
TF-IDF + PA-W2V vs TF-IDF LR	-0.0012	-0.0057 to +0.0032	0.5984
TF-IDF + PA-W2V vs TF-IDF + SGNS	+0.0049	+0.0011 to +0.0088	0.0184
PA-W2V mean LR vs SGNS mean LR	+0.0140	+0.0087 to +0.0193	<0.0001
PA-W2V mean LR vs SGNS+neg mean LR	+0.0076	+0.0029 to +0.0119	0.0007
SGNS+neg mean LR vs SGNS mean LR	+0.0064	+0.0022 to +0.0107	0.0062
DistilBERT fine-tuned vs TF-IDF LR	+0.0190	+0.0132 to +0.0249	<0.0001
DistilBERT fine-tuned vs TF-IDF + PA-W2V	+0.0202	+0.0144 to +0.0260	<0.0001

- **DistilBERT is the best classifier**: macro-F1 0.916, a difference of +0.019 over TF–IDF with a confidence interval excluding zero. Its confusion matrix is 4,524 / 448 for actual Bad and 383 / 4,519 for actual Good. It cost about 25 minutes of GPU time for five folds, against about 22 seconds for TF–IDF.
- **Each PA-W2V step improves embedding features.** SGNS 0.871 → with negation binding 0.878 → with counter-fitting 0.885, all significant. This is a real downstream gain for the Word2Vec contribution.
- **The TF–IDF hybrid adds nothing.** TF–IDF + PA-W2V (0.8955) and TF–IDF (0.8968) are statistically indistinguishable ($p = 0.60$). With bigrams available, averaged embeddings bring little

extra.

- **Domain training matters.** The untrained SST-2 model reaches 0.819, with poor calibration (ECE 0.159).
- **Continuity.** On the earlier 1K file this project's TF–IDF model reaches macro-F1 0.882, compared with 0.8715 in the earlier report. Moving to 10K raises TF–IDF to 0.897.

5.3 Original contribution: selective triage

Rather than a label for every review, the system acts automatically only when confidence reaches a threshold τ, and defers others to an analyst. For selective classification [15], the risk at threshold τ is

$$R(\tau) = \frac{\sum_i \mathbf{1}[c_i \geq \tau]\,\mathbf{1}[\hat{y}_i \neq y_i]}{\sum_i \mathbf{1}[c_i \geq \tau]}, \qquad \tau^* = \min\{\tau : R(\tau) \leq r\}, \tag{5}$$

with $c_i = \max(p_i, 1 - p_i)$. To avoid reporting a threshold chosen on the data it is tested on, τ^* for each fold is selected using the other four. A support manager can then state an error tolerance r and receive the corresponding workload.

Table 9: Selective triage with cross-fitted thresholds: the requested risk is compared with the risk achieved on unseen folds.

Model	Target risk	Coverage	Achieved risk	Deferred	Bad recall (auto)
TF-IDF LR	0.10	0.989	0.099	112	0.923
TF-IDF LR	0.05	0.841	0.049	1568	0.966
TF-IDF LR	0.03	0.711	0.029	2849	0.980
TF-IDF LR	0.02	0.613	0.020	3818	0.985
TF-IDF + PA-W2V	0.10	0.987	0.099	125	0.926
TF-IDF + PA-W2V	0.05	0.821	0.051	1766	0.962
TF-IDF + PA-W2V	0.03	0.676	0.030	3200	0.977
TF-IDF + PA-W2V	0.02	0.534	0.019	4605	0.981
DistilBERT fine-tuned	0.10	1.000	0.084	0	0.910
DistilBERT fine-tuned	0.05	0.909	0.049	895	0.947
DistilBERT fine-tuned	0.03	0.817	0.032	1808	0.967
DistilBERT fine-tuned	0.02	0.717	0.019	2798	0.982

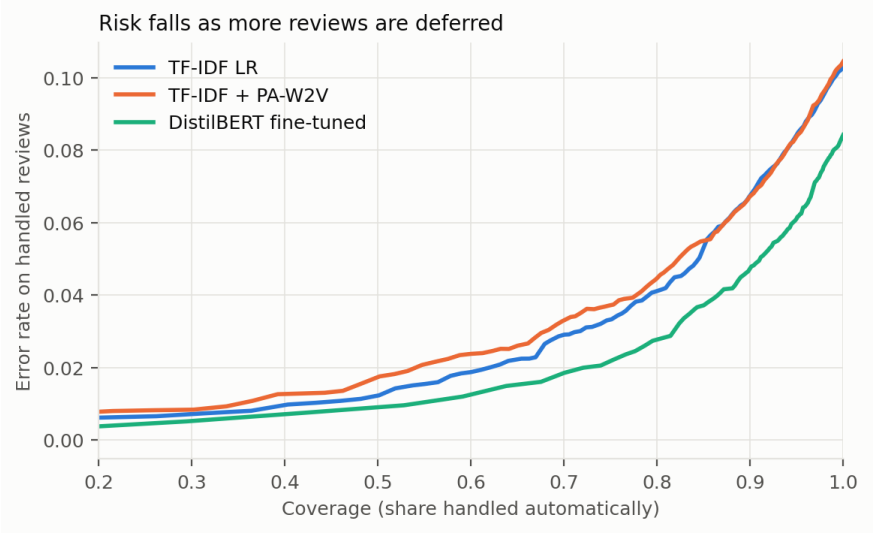


Figure 5: Risk against coverage. DistilBERT handles more reviews at every error level.

Achieved risk tracks the target closely on unseen folds. With DistilBERT at a 5% target, the system handles 91.0% of reviews with 4.93% error and sends 895 of 9,874 to an analyst. The earlier report's fixed-threshold triage has no such control.

5.4 Robustness

Table 10: Robustness: authored contrast set (30 sentences) and cross-domain test on Comcast complaints.

Model (trained on all 10K)	Contrast acc.	Comcast macro-F1	Comcast acc.	Comcast Bad recall
TF-IDF LR	0.667	0.685	0.985	0.994
TF-IDF + PA-W2V	0.733	0.675	0.984	0.993
PA-W2V mean LR	0.733	0.641	0.978	0.987
DistilBERT fine-tuned	0.867	0.699	0.982	0.988

On 30 authored minimal pairs, DistilBERT is correct on 86.7%, compared with 66.7% for TF–IDF. Embedding features help on simple negation ("the engineer was not helpful" receives $P(\text{Good})$ 0.93 under TF–IDF and 0.26 under the hybrid) but not on idioms: every model, DistilBERT included, still fails on "could not be better". This set is small and hand-written, so it is diagnostic only. On Comcast complaints Bad recall stays above 98%, but only 51 Good reviews exist, so macro-F1 of 0.64–0.70 rests on few examples. The dataset has no demographic attributes, so these tests are not a fairness assessment.

6 Case Study 3: Text Summarization

6.1 Original contribution: PAB-MMR

Maximal Marginal Relevance [16] balances relevance and redundancy. Polarity- and Aspect-Balanced MMR adds two terms. Given source sentences with vectors $\mathbf{x}_i$, polarity p_i and aspect a_i, and a target share of positive sentiment π, it greedily selects under a word budget

$$s^* = \arg\max_{i \notin S} \Big[\lambda \cos(\mathbf{x}_i, \bar{\mathbf{x}}) - (1-\lambda) \max_{j \in S} \cos(\mathbf{x}_i, \mathbf{x}_j) + \beta\, w(a_i) \mathbf{1}[a_i \notin a(S)] - \eta\, \mathrm{JS}\big(\bar{p}_{S \cup i} \,\|\, \pi\big) \Big], \quad (6)$$

where $w(a)$ is the aspect's share of the cluster and $\bar{p}$ is the word-weighted mean polarity. Jensen–Shannon divergence with base-2 logarithms is bounded in $[0,1]$, is zero exactly when the distributions match, and is symmetric, so it is a stable penalty for greedy selection. Setting $\beta = \eta = 0$ recovers standard MMR, so the method is a strict generalisation and each term can be tested by ablation.

Aspects (billing, speed, support, installation, contract) use five seed lists expanded with PA-W2V neighbours (`aspect_lexicon.json`), connecting this task to Case Study 1. Telecom polarity comes from TF–IDF models cross-fitted across folds, so a review's own label never affects its sentence scores, connecting it to Case Study 2.

6.2 Recovering sentences from unpunctuated text

The telecom corpus has no punctuation, so sentences must be recovered. To measure this with real ground truth, punctuation was stripped from 453 multi-sentence FewSum reviews and boundaries were compared with the originals.

Table 11: Sentence-boundary recovery after stripping punctuation from 453 FewSum reviews (tolerance two words).

Segmenter	Precision	Recall	F1	ROUGE-2 (test)
Cue words	0.254	0.179	0.194	0.047
Fixed 15-word windows	0.320	0.306	0.299	0.055
Punctuation restoration	**0.818**	**0.887**	**0.828**	**0.054**

The first design, a cue-word splitter, was *worse* than fixed windows and produced fragments. It was replaced with a pretrained punctuation model [20], which raises boundary F1 from 0.299 to 0.828. Fixed windows reach a similar ROUGE-2 on stripped text, as ROUGE rewards word overlap rather than readable sentences, so ROUGE alone would have missed this improvement.

6.3 Evaluation against human references

Using the FewSum gold summaries [26] and ROUGE [25], hyperparameters were selected on 40 products by ROUGE-2 and tested on 20. TextRank [24] and centroid selection serve as graph and frequency baselines. The word budget of 50 equals the mean reference length.

Table 12: FewSum test (20 products, three human references each): ROUGE F1.

Method	Words	ROUGE-1	ROUGE-2	95% CI (R-2)	ROUGE-L
Lead	48.7	0.267	0.041	0.029–0.055	0.155
Centroid	48.5	0.278	0.047	0.036–0.059	0.165
TextRank	48.3	0.281	0.045	0.036–0.055	0.164
MMR (tuned)	48.7	0.290	0.044	0.033–0.055	0.167
PAB-MMR (tuned)	48.5	0.277	0.045	0.035–0.056	0.163
PAB-MMR (fixed)	49.2	0.288	0.047	0.036–0.058	0.163
BART-large-CNN	41.9	0.265	0.043	0.029–0.061	0.163

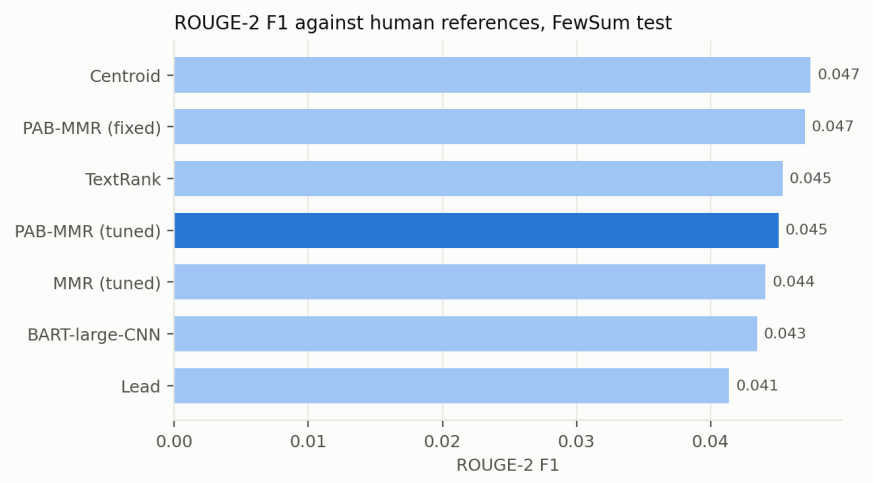


Figure 6: ROUGE-2 F1 on FewSum test. Confidence intervals overlap, so the methods are not separable.

All methods are statistically tied. PAB-MMR against MMR gives a mean ROUGE-2 difference of +0.001 ($p = 0.41$), and with 20 products this test has limited power. Scores are low across the board because human summaries paraphrase heavily and extractive methods cannot match unseen wording. The correct reading is that the balance terms *do not harm* reference similarity; they do not improve it.

6.4 Evaluation on telecom digests

No human references exist for telecom clusters, so intrinsic measures are used over 47 clusters (up to 60 reviews each) with PAB-MMR weights fixed in advance ($\lambda = 0.7, \beta = 0.3, \eta = 0.5$). *Polarity gap* is the Jensen–Shannon divergence between summary sentiment and the true share of Good reviews; *aspect recall* is the share of the cluster's segments whose aspect appears in the summary.

Table 13: Telecom digests over 47 provider and provider-by-aspect clusters (no human references).

Method	Polarity gap	Aspect recall	Redundancy	Centroid cos	Words
Lead	0.046	0.653	0.030	0.280	48.1
Centroid	0.074	0.521	0.317	0.517	48.4
TextRank	0.065	0.562	0.290	0.518	48.5
MMR	0.059	0.680	0.081	0.576	48.3
PAB-MMR	**0.019**	**0.850**	**0.072**	**0.545**	**48.6**
BART-large-CNN	0.081	–	–	–	60.1

Table 14: Paired Wilcoxon tests across the 47 telecom clusters.

Metric	Versus	PAB-MMR	Other	p
Polarity gap	MMR	0.019	0.059	<0.0001
Aspect recall	MMR	0.850	0.680	<0.0001
Redundancy	MMR	0.072	0.081	0.1334
Polarity gap	TextRank	0.019	0.065	<0.0001
Polarity gap	BART-large-CNN	0.019	0.081	0.0002

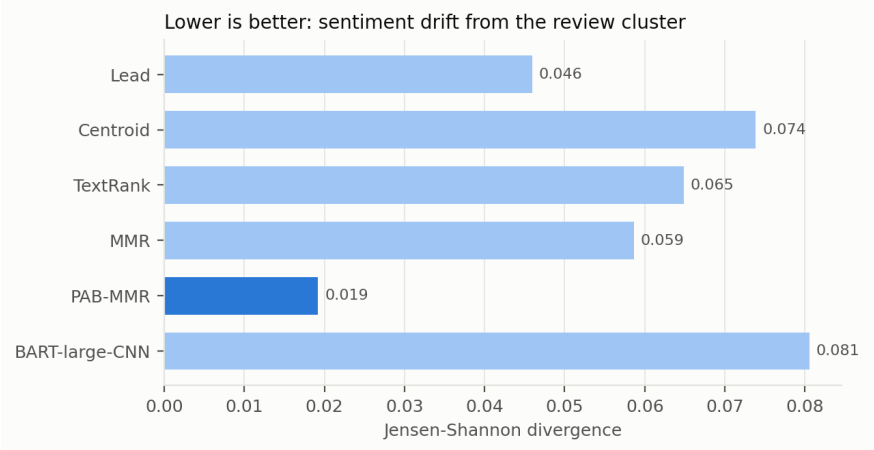


Figure 7: Sentiment drift from the source cluster. BART is evaluated with the sentence-polarity model trained on all reviews.

PAB-MMR reduces the polarity gap by 67% against MMR (0.059 to 0.019) and raises aspect recall from 0.680 to 0.850, both significant. The cost is a small fall in centroid similarity (0.576 to 0.545). **An important caveat:** polarity gap and aspect recall are what PAB-MMR explicitly optimises, so this shows that the method does what it is designed to do, not that human readers prefer it. A reader study would be needed for that claim.

Table 15: Generated digests (lowercase source text, segments separated by bars).

Cluster	Summary
virgin PAB-MMR	i had been a loyal virgin media customer for years | i have been with virgin broadband for a few months now and cant wait for my contract to end | i have found their internet service and customer service superb | when virgin media works it can be a very good service
virgin MMR	i had been a loyal virgin media customer for years | i have been with virgin broadband for a few months now and cant wait for my contract to end | at present the service is not intermittent | when virgin media works it can be a very good service
bt PAB-MMR	so another call to bt i made | i have been a customer of zen internet for ten years | later i was told they will get the original phone number back once the phone line is restored | really happy with the service and product received from bt shop
bt MMR	so another call to bt i made | later i was told they will get the original phone number back once the phone line is restored | i have been a customer of zen internet for ten years | really happy with the service and product received from bt shop

BART-large-CNN [17], a news model, tends to copy from the start of its truncated input (novel token rate 2.4%) and in one cluster produced text about other providers, as the concatenated input mixes mentions. Extractive digests remain directly traceable to source reviews, which supports human verification. The digests also show remaining weaknesses: sentences can be off-topic, and provider clusters are found by keyword, so a review mentioning a competitor can enter the wrong cluster.

7 Case Study 4: Factoid Question Answering

7.1 Pipeline

The system follows the classic retriever–reader architecture for factoid QA [28]:

1. **Question processing.** Rules map a question to an expected type: MONEY, DURATION, COUNT, TIME, DATE or ENTITY.
2. **Passage retrieval.** Complaints are split into three-sentence overlapping passages, indexed by BM25 [18], with dense and fused alternatives.
3. **Answer extraction.** A SQuAD-trained reader [19,22] scores spans in the top ten passages.
4. **Answer-type gate (original).** Spans that do not fit the expected type are discarded, so a MONEY question cannot return “Modem”.
5. **Calibrated abstention (original).** Each candidate receives

$$\text{score} = (s_{\text{start}} + s_{\text{end}} - s_{\text{CLS}}) + \frac{w}{\text{rank}}, \tag{7}$$

the reader’s advantage over its no-answer option plus a retrieval bonus. If the best score falls below a threshold τ tuned on development questions, the system answers “not found in the corpus”. Every answer is returned with its source complaint.

7.2 Why a new gold set was needed

No public factoid set exists for telecom complaints. A set of 132 questions was therefore written against seeded samples of complaints (`qa_gold.jsonl`). Every answer was checked programmatically to occur verbatim in its source complaint. Thirty questions ask for facts the corpus should not contain. Questions were split into 44 development and 88 test before any system was run.

7.3 Public benchmark

Table 16: Reader benchmark on a seeded 2,000-question SQuAD v1.1 dev sample with the gold paragraph supplied.

Reader	EM	95% CI	F1	Seconds
DistilBERT-SQuAD	0.785	0.767–0.802	0.860	25
RoBERTa-SQuAD2	0.844	0.829–0.860	0.907	43

Both readers perform in a credible range on SQuAD [21], which confirms correct integration before domain evaluation. RoBERTa scores higher here, and development selection on telecom questions chose it as well.

7.4 Retrieval

Table 17: Retrieval of the evidence complaint among 5,629 documents (held-out answerable questions).

Retriever	Recall@1	Recall@5	Recall@20	MRR
BM25	0.603	0.824	0.868	0.699
SGNS dense	0.162	0.324	0.500	0.234
PA-W2V dense	0.132	0.265	0.397	0.192
BM25 + SGNS	0.250	0.618	0.853	0.427
BM25 + PA-W2V	0.235	0.559	0.853	0.391
BM25 rerank SGNS 0.25	0.603	0.824	0.868	0.703
BM25 rerank SGNS 0.5	0.603	0.824	0.868	0.700
BM25 rerank PA-W2V 0.25	0.603	0.838	0.882	0.697
BM25 rerank PA-W2V 0.5	0.588	0.838	0.882	0.695

This is a negative result for embeddings. Dense retrieval is far weaker than BM25, and equal-weight reciprocal rank fusion [23] damages BM25 (MRR 0.699 to about 0.40), because factoid questions depend on exact numbers and rare names, which static vectors encode poorly (all numbers become one token). Gentle reranking of BM25's top 100 neither helps nor harms measurably. Word2Vec suits sentiment and aspect expansion here, but not precise fact retrieval. Questions written while reading a complaint also share its wording, which favours BM25 (Section 10).

7.5 End-to-end results

Configuration was selected on development questions: BM25 reranked with SGNS at weight 0.25, RoBERTa, gate on, $w = 2$, $\tau = 3.27$.

Table 18: Ablation on 68 held-out answerable questions.

System	EM	95% CI	F1
Rule extractor (regex types)	0.324	0.221–0.441	0.338
BM25 + RoBERTa-SQuAD2, no gate	0.588	0.471–0.706	0.632
BM25 + RoBERTa-SQuAD2, gate	0.618	0.500–0.735	0.642
BM25+SGNS + RoBERTa-SQuAD2, gate	0.544	0.426–0.662	0.561
BM25+PA-W2V + RoBERTa-SQuAD2, gate	0.500	0.382–0.618	0.532
BM25 rerank PA-W2V 0.25 + RoBERTa-SQuAD2, gate	0.632	0.515–0.735	0.650
Dev-selected retriever + DistilBERT-SQuAD, gate	0.574	0.456–0.677	0.596
Dev-selected system	**0.618**	**0.500–0.735**	**0.642**

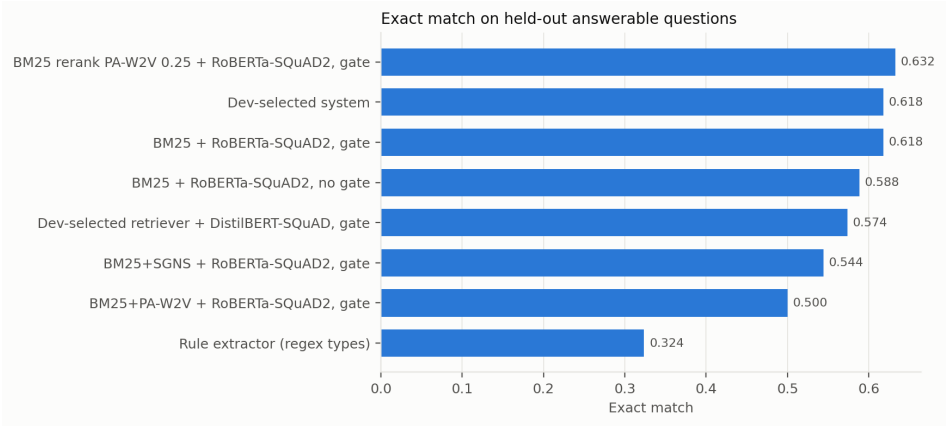


Figure 8: Exact match by system on held-out answerable questions.

A reader clearly beats regular-expression extraction (0.618 against 0.324, with non-overlapping intervals). The gate adds about three points (0.588 to 0.618), but with 68 questions the intervals are wide (± 0.12), so that gain and the differences among reader systems are suggestive only.

Table 19: Abstention on all 88 held-out questions (68 answerable, 20 unanswerable).

Policy	Overall acc.	Coverage	Answered EM	Answered F1	Rejected
Always answer	0.477	1.000	0.618	0.642	0.000
Abstain below $\tau = 3.27$	**0.602**	**0.721**	**0.755**	**0.769**	**0.800**

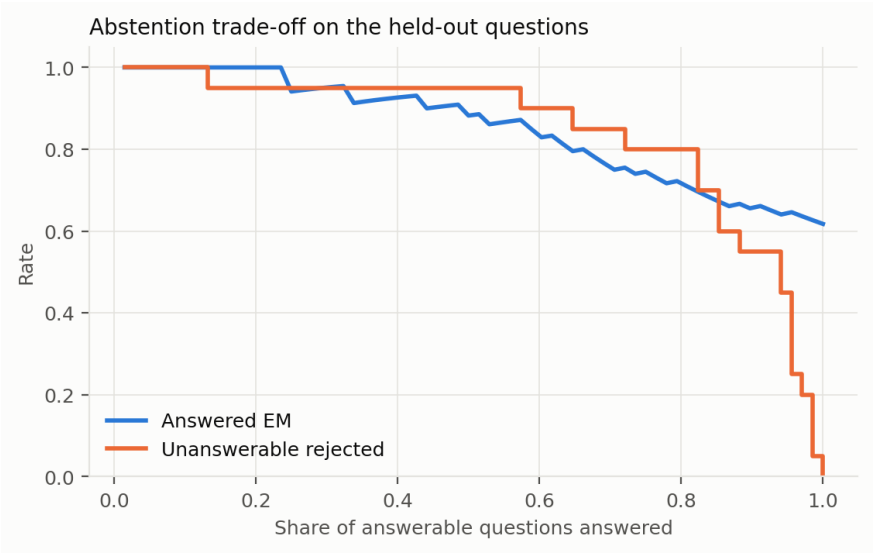


Figure 9: Lowering the threshold answers more questions but rejects fewer unanswerable ones.

Without abstention, the system invents an answer for every unanswerable question. With the threshold chosen on development data, it rejects 80% of them, and answered exact match rises from 0.618 to 0.755, while still answering 72% of answerable questions. Overall accuracy rises from 0.477 to 0.602. This is the largest practical gain in the QA case study.

Table 20: Dev-selected system by detected answer type.

Type	Questions	EM	F1
COUNT	1	1.000	1.000
DATE	2	0.500	0.500
DURATION	13	0.615	0.692
ENTITY	24	0.625	0.625
MONEY	26	0.577	0.603
TIME	2	1.000	1.000

7.6 Error analysis

Inspection of errors (`test_errors.csv`) shows three main causes:

- **Wrong complaint retrieved.** Many complaints share the same pattern, so "How much in lost wages…" can find a similar case with another figure. A plausible amount from the wrong complaint fits the expected type and passes the gate.
- **Span boundaries.** "158.00 a month" against gold "158.00" is correct in substance but counted wrong by exact match; token F1 captures this partly.
- **Reader confusion on complaint language.** "Which provider…" can return Comcast itself, since the reader lacks the idea of "another provider".

8 Integrated System

All four components run together in `05_integrated_system.ipynb`, which also runs the unit tests (`tests/test_core.py`). Outputs are recorded in `demo.json`.

Table 21: Integrated system: live triage output.

Incoming review	Prediction	P(Good)	Decision
my broadband keeps disconnecting and customer service has not fixed it	Bad	0.039	automatic
the engineer arrived on time and the router was set up in minutes	Good	0.939	automatic
it was fine i suppose though the price went up after a few months	Good	0.861	automatic
the speed is good but nobody answers when you call and the bill was wrong once	Good	0.800	automatic
great deal at first however the router needed replacing twice and they sorted it eventually	Good	0.982	automatic

Table 22: Integrated system: live question answering with evidence or abstention.

Question	Answer	Score	Complaint
What was the cancellation fee for the customer who moved to a neighborhood not serviced by Comcast?	$756.82	11.86	2562
Which insurance company refused to pay the full electrician estimate after faulty cable wiring?	Liberty Mutual	10.58	4307
How much does Comcast charge for 5G home internet in Seattle?	(abstained)	2.91	–

Digest for Virgin reviews: *im not happy with virgin media | i had phone call to virgin mobile customer services just to change my but i had silly conversation and no have been virgin broadband around more*

than years and last year i decided to change my mobile to virgin and now i am realy p off with your | i want my contract cancelled

The integrated system uses TF–IDF for live triage because it is fast and explainable; DistilBERT remains the most accurate option where a GPU is available. The demonstration exposes a limitation: the mixed review "the speed is good but nobody answers when you call" is confidently labelled Good and handled automatically, so deferral did not catch it. Deferral lowers error in aggregate (Table 9) but does not guarantee catching any individual review.

9 Comparison with Supplied References

Reported scores from different datasets cannot be ranked directly because labels, sample sizes and protocols differ.

Table 23: Critical comparison with the four supplied papers.

Reference	Focus and reported contribution	Relationship to this project
[1] Abdelmotaleb et al. (2026)	Compares TF–IDF, Word2Vec, FastText, Doc2Vec, BERT, classifiers and PCA pooling for telecom star ratings, with computational costs.	Source of the dataset and of the accuracy-versus-cost framing. This project uses the 10K file released with it, not the 367,000-review corpus. It adds selective deferral, summarization and QA.
[2] Siswa et al. (2026)	Compares TF–IDF, Word2Vec and FastText with classical classifiers on 3,441 Indonesian comments; reports 85.2% for SVM + FastText.	Here FastText was the weakest embedding feature (0.857), which reinforces that representation choice depends on domain and language.
[3] Joseph and Odumuyiwa (2026)	Surveys text representations from classical vectors to sentence embeddings.	Provides the conceptual framing. The pooling result here (Table 5) illustrates its point about matching representation granularity to the task.
[4] Khare and Pandey (2026)	Reviews Word2Vec, Bi-LSTM and transformer department classification in healthcare, with overfitting concerns.	Motivates this project's emphasis on leakage control, cross-fitting and held-out evaluation.

10 Threats to Validity and Disclosed Limitations

- **Single-annotator QA gold.** One annotator wrote the 132 questions, reading the complaints directly, so they share vocabulary with their sources. This favours lexical retrieval, and the 88 test questions give wide confidence intervals.
- **Changes after seeing test results.** These are disclosed in full: (i) the answer-type pattern was broadened after a test error revealed a bug, leaving the dev-selected exact match unchanged at 0.618; (ii) reranking modes were added after a first run in which fusion performed badly, though selection remained on development data; (iii) the cue-word segmenter was replaced after its measured failure on the segmentation benchmark; (iv) after the first summarization run, BART's telecom length limit was lowered towards the word budget and the minimum telecom segment length was raised from four to five words. The QA test set is therefore not strictly blind.
- **Self-aligned metrics.** PAB-MMR's telecom advantage is on metrics it optimises. No reader study has been conducted.
- **Small benchmarks.** FewSum test has 20 products; the contrast set has 30 sentences; Comcast has 51 Good reviews.
- **Embedding held-out pairs are curated.** They were chosen by hand from common telecom words. A different list would change the numbers.

- **Keyword-based clusters.** Provider and aspect assignment is lexical and has not been separately validated.
- **Labels are proxies.** Star ratings do not always reflect the text, and the data comes from one review source per corpus, so results may not generalise across providers or countries.
- **No operational evidence.** No claim is made about reduced waiting time or cost; a field evaluation would be required.

11 Conclusion and Future Work

This case study delivers all four assigned topics as one working, evaluated system. Its main findings are:

1. **Word2Vec.** Plain embeddings confuse opposites. PA-W2V cut held-out antonym similarity from 0.683 to 0.373 and improved embedding-based sentiment significantly, though sentence-level pooling hides much of the gain and SimLex shows no change.
2. **Sentiment.** Fine-tuned DistilBERT achieved macro-F1 0.916. Cross-fitted selective triage met its 5% risk target on unseen folds while automating 91% of reviews.
3. **Summarization.** PAB-MMR reduced sentiment drift by 67% and increased aspect recall without measurable harm to ROUGE. Punctuation restoration made unpunctuated reviews usable, with boundary F1 of 0.83.
4. **Factoid QA.** Answer gating and abstention rejected 80% of unanswerable questions and raised overall accuracy from 0.477 to 0.602. BM25 outperformed embeddings for retrieval, which is a useful negative result.

The main lesson is that each method should be chosen by evidence for each task. Embeddings improved sentiment representation and failed at fact retrieval, and the most valuable feature across tasks was knowing when not to decide.

Next steps are to obtain independent annotation of the QA gold set with agreement statistics, run a reader study on digests, use contextual embeddings for sentence-level negation, replace keyword clusters with validated aspect models, and evaluate the triage policy with support analysts before making any operational claim.

References

[1] H. Abdelmotaleb, C. Mcneile, and M. Wojtyś, "A comparative study of word embedding techniques for classification of star ratings," *Expert Systems with Applications*, vol. 297, article 129037, 2026. doi: 10.1016/j.eswa.2025.129037.

[2] T. A. Y. Siswa, M. Fattah, M. A. Fahrozi, D. F. Rahman, and F. I. Ramadhani, "Optimizing Sentiment Classification on IKN Development: A Comparative Study of TF-IDF, Word2Vec, and FastText Embeddings," *Engineering Proceedings*, vol. 137, no. 20, 2026. doi: 10.3390/engproc2026137020.

[3] P. K. Joseph and V. Odumuyiwa, "From word to sentence embedding and beyond: Bridging the gap in text representation survey," *Natural Language Processing Journal*, vol. 15, article 100214, 2026. doi: 10.1016/j.nlp.2026.100214.

[4] H. Khare and D. Pandey, "Deep Neural Architectures for Automated Department Classification: A Review of Word2Vec, Bi-LSTM, and Transformer-Based Models in Healthcare," in *2026 IEEE 6th IC2PCT*, 2026. doi: 10.1109/IC2PCT68894.2026.11584372.

[5] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," arXiv:1301.3781, 2013.

[6] T. Mikolov, I. Sutskever, K. Chen, G. Corrado, and J. Dean, "Distributed representations of words and phrases and their compositionality," in *NeurIPS*, 2013.

[7] P. Bojanowski, E. Grave, A. Joulin, and T. Mikolov, "Enriching word vectors with subword information," *TACL*, vol. 5, pp. 135–146, 2017.

[8] N. Mrkšić et al., "Counter-fitting word vectors to linguistic constraints," in *NAACL-HLT*, 2016.

[9] M. Faruqui et al., "Retrofitting word vectors to semantic lexicons," in *NAACL-HLT*, 2015.

[10] B. L. Monroe, M. P. Colaresi, and K. M. Quinn, "Fightin' words: Lexical feature selection and evaluation for identifying the content of political conflict," *Political Analysis*, vol. 16, no. 4, 2008.

[11] B. Pang, L. Lee, and S. Vaithyanathan, "Thumbs up? Sentiment classification using machine learning techniques," in *EMNLP*, 2002.

[12] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, "DistilBERT, a distilled version of BERT," arXiv:1910.01108, 2019.

[13] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *ICML*, 2017.

[14] J. Platt, "Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods," in *Advances in Large Margin Classifiers*, 1999.

[15] Y. Geifman and R. El-Yaniv, "Selective classification for deep neural networks," in *NeurIPS*, 2017.

[16] J. Carbonell and J. Goldstein, "The use of MMR, diversity-based reranking for reordering documents and producing summaries," in *SIGIR*, 1998.

[17] M. Lewis et al., "BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension," in *ACL*, 2020.

[18] S. Robertson and H. Zaragoza, "The probabilistic relevance framework: BM25 and beyond," *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, 2009.

[19] Y. Liu et al., "RoBERTa: A robustly optimized BERT pretraining approach," arXiv:1907.11692, 2019.

[20] O. Guhr, A.-K. Schumann, F. Bahrmann, and H. J. Böhme, "FullStop: Multilingual punctuation prediction," in *SwissText*, 2021.

[21] P. Rajpurkar, J. Zhang, K. Lopyrev, and P. Liang, "SQuAD: 100,000+ questions for machine comprehension of text," in *EMNLP*, 2016.

[22] P. Rajpurkar, R. Jia, and P. Liang, "Know what you don't know: Unanswerable questions for SQuAD," in *ACL*, 2018.

[23] G. V. Cormack, C. L. A. Clarke, and S. Büttcher, "Reciprocal rank fusion outperforms Condorcet and individual rank learning methods," in *SIGIR*, 2009.

[24] R. Mihalcea and P. Tarau, "TextRank: Bringing order into text," in *EMNLP*, 2004.

[25] C.-Y. Lin, "ROUGE: A package for automatic evaluation of summaries," in *ACL Workshop on Text Summarization Branches Out*, 2004.

[26] A. Bražinskas, M. Lapata, and I. Titov, "Few-shot learning for opinion summarization," in *EMNLP*, 2020.

[27] F. Hill, R. Reichart, and A. Korhonen, "SimLex-999: Evaluating semantic models with (genuine) similarity estimation," *Computational Linguistics*, vol. 41, no. 4, 2015.

[28] D. Jurafsky and J. H. Martin, *Speech and Language Processing*, 3rd ed. draft.

A Evidence Index

Each headline claim maps to a file in the public repository.

Claim	Evidence file
Word2Vec metrics, five folds	results/word2vec/intrinsic_by_fold.csv
Counter-fitting grid	results/word2vec/counterfit_tuning.csv
Sentiment scores and intervals	results/sentiment/scores.csv
Paired sentiment tests	results/sentiment/comparisons.csv
Per-review out-of-fold probabilities	results/sentiment/oof_probabilities.csv
Selective triage	results/sentiment/selective.csv
ROUGE per product	results/summarization/fewsum_by_group.csv
Boundary recovery	results/summarization/segmentation.csv
Telecom digest metrics, 47 clusters	results/summarization/telecom_by_cluster.csv
Generated digests	results/summarization/telecom_digests.csv
SQuAD benchmark	results/qa/squad_reader_benchmark.csv
QA configuration search on development data	results/qa/dev_search.csv
Every held-out QA prediction	results/qa/test_predictions.jsonl
Abstention results	results/qa/test_selective.json
Integrated outputs and unit tests	results/integrated/demo.json, pytest.txt

B Execution Record

Each notebook's executed run, with its code, cell outputs and log, is public on Kaggle. Each run cloned this repository from GitHub, and the results in this report come from these runs.

Case study	Public Kaggle notebook
Word2Vec	https://www.kaggle.com/code/bharatchandra1234/telelens-01-word2vec
Sentiment Classification	https://www.kaggle.com/code/bharatchandra1234/telelens-02-sentiment
Text Summarization	https://www.kaggle.com/code/bharatchandra1234/telelens-03-summarization
Factoid QA	https://www.kaggle.com/code/bharatchandra1234/telelens-04-factoid-qa
Integrated system	https://www.kaggle.com/code/bharatchandra1234/telelens-05-integrated

Excerpts from the run logs and unit test output:

```
device cuda Tesla T4
9874 reviews 617573 tokens 11434 types
5629 documents 34054 passages 44 dev 88 test 68 answerable test
{'retriever': 'BM25 rerank SGNS 0.25', 'reader': 'RoBERTa-SQuAD2', 'gate': True, 'weight': 2.0}
tau 3.2721939086914062
14 passed (pytest, tests/test_core.py)
```

C Reproducibility

1. Clone https://github.com/BharatChandra-sys/NLP-case-study.
2. Open a notebook from notebooks on Kaggle, enable a GPU and internet, and attach the Comcast dataset.

3. Run all cells. The first cell clones the repository and downloads public data; results are written to `results/`.
4. Seeds are fixed at 42. Tables in this report are generated from `results/` by `make_tables.py`.

Reproducibility check. A second complete run from a fresh GitHub clone reproduced the from-scratch SGNS and PA-W2V metrics and the TF–IDF, DistilBERT, summarization and QA headline scores. Two small differences were observed. The gensim baselines train on several threads and are not exactly repeatable: their metrics varied by up to 0.038, and FastText sentiment macro-F1 moved from 0.863 to 0.857. DistilBERT's 5% deferral count moved from 892 to 895 reviews. No conclusion in this report depends on these differences.

D Selected Implementation Fragment

```
def answer(self, question, gate=True, retrieval_weight=2.0):
    ranked, top, contexts, all_spans = self.read(question)
    kind = answer_type(question)
    best = {"answer": "", "score": -1e9, "doc_id": None, "evidence": "", "type": kind}
    for position, spans in enumerate(all_spans):
        for span in spans:
            if gate and not passes_gate(kind, span["text"]):
                continue
            score = span["margin"] + retrieval_weight / (position + 1)
            if score > best["score"]:
                best = {"answer": span["text"], "score": score,
                        "doc_id": self.retriever.passages[top[position]]["doc_id"],
                        "evidence": contexts[position], "type": kind}
    return best, ranked
```